

CPDS Laboratory Hands-On

OpenMP: Brief tutorial

J.R. Herrero

Fall 2023



UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH

Index

Index	1
1 A very practical introduction to OpenMP (I)	2
1.1 Computing number Pi	2
1.2 Parallelisation with OpenMP	3
1.2.1 Defining the parallel region and using implicit tasks	3
1.2.2 Using synchronisation to avoid data races	3
1.2.3 Summary of code versions	4
1.3 Test your understanding	4
1.4 Observing overheads	5
1.4.1 Synchronisation overheads	5
2 A very practical introduction to OpenMP (Part II)	6
2.1 Parallelisation with OpenMP	6
2.1.1 Tasking execution model and use of explicit tasks	6
2.1.2 Summary of code versions	7
2.2 Test your understanding	8
2.3 Observing overheads	8
2.3.1 Thread creation and termination	8
2.3.2 Task creation and synchronisation	8
3 A very practical introduction to OpenMP (Part III)	10
3.1 Work-sharing constructs	10
3.2 Additional examples	10

Note:

- All files necessary to do this laboratory hands-on are available in a compressed tar file available from the following location: `/scratch/nas/1/cpds0/sessions/H0_OpenMP.tar.gz`. Copy it to your home directory in `boada.ac.upc.edu` and uncompress it with this command line: `"tar -zxvf H0_OpenMP.tar.gz"`.

Examples 1

A very practical introduction to OpenMP (I)

This laboratory assignment has been prepared with the purpose of introducing the main constructs in the OpenMP extensions to the C programming language. In each of the two sessions you will first go through a set of different code versions (some of them not correct) for the computation of number Pi in parallel; and then you will be presented with a set of simple examples that will be helpful to practice the main components of the OpenMP programming model being introduced. We ask you to fill in the questionnaire in the deliverable part for this second laboratory assignment. The session will finish observing the overheads introduced by the use of different synchronisation constructs in OpenMP.

1.1 Computing number Pi

In the documentation for the previous laboratory assignment we already showed how to compute the number Pi using numerical integration. To distribute the work for the parallel version each processor will be responsible for executing some iterations of the loop `i` in Figure 1.1. The parallelisation should also guarantee that there are no data races when accessing to variable `sum` in order to accumulate the contribution of each iteration.

```
static long num_steps = 100000;
void main () {
    double x, pi, step, sum = 0.0;
    step = 1.0/(double) num_steps;

    for (long int i=0; i<num_steps; ++i) {
        x = (i+0.5)*step;
        sum += 4.0/(1.0+x*x);
    }
    pi = step * sum;
}
```

Figure 1.1: Pi computation

In order to parallelise the sequential code we will proceed through a set of versions "`pi-vx.c`" inside the `HO.OpenMP/pi` directory, being `x` the version number. We provide two entries in the `Makefile` to compile them: "`make pi-vx-debug`" and "`make pi-vx-omp`". The binary generated with the "`make pi-vx-debug`" option prints which iterations are executed by each thread and the value of Pi that is computed, which will be useful to understand what the program is doing. You can use the script `run-debug.sh` to interactively execute the binary with a very small input value (by default set to 32 iterations in the script). In order to time the parallel execution you need to choose the second option "`make pi-vx-omp`" and queue the execution of the `submit-omp.sh` script, which is using a much larger input value (by default set to 100000000 iterations in the script). In both cases, after the script name you only have to specify `pi-vx`. In all these scripts, the number of processors to be used is by default set

to 4. You can change both the number of iterations and the number of processors by adding two extra arguments when executing or submitting for execution the script (e.g. `sbatch submit-omp.sh pi-vx 1000000000 8`).

1.2 Parallelisation with OpenMP

In this section you will first learn how to define a parallel region to create a "team of threads" that will execute in parallel. Each one of the threads created in the parallel region will execute a so-called "implicit task" that corresponds with the body of the parallel region. You will also see how to split the work among these implicit tasks (as many as threads in the parallel region) and how to guarantee the correct access to private and shared variables (using different synchronisation constructs). Then in the next session you will go one step further into the so-called "explicit tasks", a much more versatile way to express parallelism in OpenMP. First, one of the threads (implicit task) in the team will be selected as the responsible for generating the explicit tasks that will be executed by the other threads (implicit tasks) in the team, learning the simplest way to synchronise these explicit tasks through task barriers. In subsequent sessions you will dive deeper into nested explicit tasks to see how powerful the tasking model in OpenMP is.

1.2.1 Defining the parallel region and using implicit tasks

1. Compile and run the initial sequential code `pi-v0.c`. This initial version introduces the use of `omp_get_wtime` runtime call to measure wall clock execution time. Execute both the `-debug` and `-omp` binaries to have a reference for the result that is obtained and the execution time for the sequential version. Note that the Pi value result may be different for the `run-debug.sh` (with only 32 iterations) and `submit-omp.sh` (with 1000000000 iterations) executions.
2. In a first attempt to parallelise the sequential code, `pi-v1.c` introduces the `parallel` construct, which creates/activates the team of threads. With `parallel` all threads in the team created/activated replicate the execution of the body of the parallel regions delimited by the open/close curly brackets. However, just adding `parallel` makes our parallel code incorrect, as you will realise when executing the `-debug` binary; in order to understand why, we need to know that in OpenMP all variables used inside the parallel region are shared by default unless declared private (either by defining the variable inside the scope of the parallel region or declaring it `private` in the `parallel` construct). In particular, in this code the (shared) access to the loop control variable `i` and the temporary variable `x` causes the data races that make the execution not correct.
3. In order to partially correct it, `pi-v2.c` adds the `private` clause for variables `i` and `x`. By executing the `-debug` binary you will notice that when `i` is private each thread executes all iterations of the loop, so we are not taking benefit of parallelism.
4. In order to avoid the total replication of work in the previous version, `pi-v3.c` uses the runtime call `omp_get_num_threads()` and changes the for loop in order to distribute the iterations of the loop among the participating threads. Compile and execute the `-debug` binary. Which iterations is each thread executing? Notice that the result is not correct due to a race condition (perhaps you don't see the error when using the small input value, but for sure you will see if you execute the `-omp` binary). The access to which variable is causing the data race, and therefore the incorrect result?

1.2.2 Using synchronisation to avoid data races

1. Next version `pi-v4.c` uses the `critical` construct to provide a region of mutual exclusion, that is a region of code in which only one thread executes it at any given time. Compile and execute the `-debug` binary to check the result. This version should be correct although, as you will realise when submitting the execution of the `-omp` binary, it introduces large synchronisation overheads.
2. Version `pi-v5.c` uses the `atomic` construct to guarantee the indivisible execution of read-operate-write operations, which is more efficient than the `critical` construct used in the previous version.

Compare the execution time of both `pi-v4` and `pi-v5` when submitting the `-omp` binaries for execution. In any case, still the execution of `pi-v5.c` is very inefficient.

3. In order to reduce the amount of synchronisation needed to protect the access to `sum`, version `pi-v6.c` defines a private copy for each thread `sumlocal` that is used to accumulate its partial contribution to the global variable `sum`. This new variable is initialised to 0 and before the thread finishes it accumulates its value into `sum` making use of either `atomic` or `critical`. Check that the result is correct by executing the `-debug` binary and check scalability by submitting the `-omp` binary.
4. Since the previous code transformation is very common, **OpenMP** offers the `reduction` clause. Version `pi-v7.c` makes use of this `reduction` clause. The compiler creates a private copy of the reduction variable which is used to store the partial result computed by each thread; this variable is initialised to the neutral value of the operator specified (0 in the case of `+`). At the end of the region, the compiler ensures that the shared variable is properly updated combining the partial results computed by each thread using the operator specified (`+` in this case). Run it and compare the execution time of `pi-v7` with the previous three versions in terms of execution time.
5. Finally `pi-v8.c` introduces the `barrier` synchronisation construct, as a way to define a point where all threads wait for each other to arrive. Execute the `-omp` binary and notice the partial values for `pi` that each thread is printing after the thread `barrier`, once its computation loop. Where is the reduction operation specified in the `parallel` construct happening?

1.2.3 Summary of code versions

The following table summarises all the codes used during the process. Changes accumulate from one version to the following.

Version	Description of changes	Correct?
v0	Sequential code. Makes use of <code>omp_get_wtime</code> to measure execution time	yes
v1	Added <code>parallel</code> construct and <code>omp_get_thread_num()</code>	no
v2	Added <code>private</code> for variables <code>x</code> and <code>i</code>	no
v3	Manual distribution of iterations using <code>omp_get_num_threads()</code>	no
v4	Added <code>critical</code> construct to protect <code>sum</code>	yes
v5	Added <code>atomic</code> construct to protect <code>sum</code>	yes
v6	Private variable <code>sumlocal</code> and final accumulation on <code>sum</code>	yes
v7	Use of <code>reduction</code> clause on <code>sum</code>	yes
v8	Use of <code>barrier</code> construct	yes

1.3 Test your understanding

Next you will go through a set of simple examples that will be helpful to practice the components of the OpenMP programming model that have been used to parallelise the Pi computation. For each example answer the question(s) in the questionnaire that you will deliver as part of the deliverable for this second laboratory assignment. In order to follow them, you will need:

- The set of files inside the directory `H0_OpenMP/openmp/Part1`. Open each file in the directory (the examples are ordered), compile and execute; then answer the question(s) in the questionnaire associated to it. If asked in the program, do the required modifications.
- In case you want (or need) to know more details about the **OpenMP** constructs, we suggest you either use the complete specification document or the short reference guide that you will find either at the www.openmp.org website or at Atenea.

Use the appropriate entry in the `Makefile` to individually compile each program (e.g. `"make 1.hello"`). In order to do the executions simply run each program interactively (e.g. `"./1.hello"` or `"OMP_NUM_THREADS=4 ./1.hello"` if you want to externally set the number of threads to be used).

1.4 Observing overheads

1.4.1 Synchronisation overheads

To finish this session we ask you to compile and execute four different versions of the Pi computation parallel program, each one making use of a different synchronisation mechanism to perform the update of the global variable `sum`. Codes are available inside directory `H0_OpenMP/overheads`.

- `pi_omp_critical.c`: a `critical` region is used to protect every access to `sum`, ensuring exclusive access to it. This version is equivalent to `pi-v4`.
- `pi_omp_atomic.c`: it makes use of `atomic` to guarantee atomic (indivisible) access to the memory location where variable `sum` is stored. This version is equivalent to `pi-v5`.
- `pi_omp_sumlocal.c`: a “per-thread” private copy `sumlocal` is used followed by a global update at the end using only one `critical` region. This version is equivalent to `pi-v6`.
- `pi_omp_reduction.c`: it makes use of the `reduction` clause applied to the global variable `sum`. This version is equivalent to `pi-v7`.

Take a look at the four different versions and make sure you understand them. How many synchronisation operations (`critical` or `atomic`) are executed in each version? Notice that function `difference` returns the difference in time between the ideal parallel execution time (sequential divided by number of threads used) and the real parallel execution when using each kind of synchronisation. This value is printed at the end of the execution.

Compile all four versions (use the appropriate entries in the `Makefile`) and queue the execution of the binaries generated using the `submit-omp.sh` script (which requires the name of the binary, the number of iterations for Pi computation and the number of threads). Answer the following questions:

1. If executed with only 1 thread and 100.000.000 iterations, do you notice any major overhead in the execution time caused by the use of the different synchronisation mechanisms? You can compare with the baseline execution time of the sequential version in `pi_sequential.c`.
2. If executed with 4 and 8 threads and the same number of iterations, do the 4 programs benefit from the use of several processors in the same way? Can you guess the reason for this behaviour?

Take note of all the results that you obtain and reach your conclusions about the overheads associated with these `OpenMP` constructs. Can you quantify (in microseconds) the cost of each individual synchronisation operation (`critical` or `atomic`) that is used?. You will have to write your conclusions in the appropriate section of the deliverable for this second laboratory assignment.

Examples 2

A very practical introduction to OpenMP (Part II)

As in the previous session, in this one you will first go through a set of different code versions (some of them not correct) for the computation of number Pi in parallel but now using the tasking model. After that you will be presented with a set of simple examples that will be helpful to practice the main components of the OpenMP programming model being presented. We ask you to fill in the rest of the questionnaire in the deliverable for this second laboratory assignment. The session will finish observing the overheads related with the creation of parallel regions and tasks in OpenMP.

2.1 Parallelisation with OpenMP

In this section you will go one step further into the so-called "explicit tasks", a much more versatile way to express parallelism in OpenMP. First, one of the threads (implicit task) in the team will be selected as the responsible for generating the explicit tasks that will be executed by the other threads (implicit tasks) in the team, learning the simplest way to synchronise these explicit tasks through task barriers. In subsequent sessions you will dive deeper into nested explicit tasks to see how powerful the tasking model in OpenMP is.

As you did in the previous session, use the two entries provided in the `Makefile` in order to compile the different versions that will be explored in this session: `"make pi-vx-debug"` and `"make pi-vx-omp"`. Execute the binaries generated using the `run-debug.sh` (interactive execution) and `submit-omp.sh` (execution queue) scripts.

2.1.1 Tasking execution model and use of explicit tasks

1. In `pi-v9.c` we have changed the code so that the original loop is manually divided in two halves, each one computing half of the loop iterations. For each loop then an explicit task is defined, using the `task` construct. When a thread encounters a `task` construct it generates a task and places it in a pool of tasks; any thread in the team can execute it. In other words, the `task` construct provides a way of defining a deferred unit of computation that can be executed by any thread in the team of threads. Notice that each task defines a private copy of variables `i` and `x`. If you execute the `-debug` binary you will see two things: 1) each iteration of the loop is executed 4 times; and of course, 2) the value of pi is not correct. Can you guess why these two effects are happening? In order to give an answer to 1) think about which thread(s) is(are) generating tasks and how many explicit tasks are generated in total.
2. In order to solve problem 1) above, `pi-v10.c` makes use of the `single` construct so that only one thread generates the tasks; the others threads will skip the `single` construct and wait at its end; in the meanwhile, they will look in the pool of tasks for tasks to execute. Compile and run the `-debug` binary for this version and check the result that is obtained. Why the result is still not correct?

3. In order to have the correct value for pi we need to do a couple of things in our program. First, the thread that generates the two tasks has to wait for them to complete before executing the statement "pi = step * sum"; notice that this statement needs a correct value for variable `sum`, which is jointly computed by the two tasks being generated. This will require task synchronisation. And second, the two tasks have to contribute with their partial result for variable `sum`; this will require the use of one of the constructs used in versions v4 (critical), v5 (atomic) or v7 (reduction). Assuming that you would suggest the use of reduction (it was the most efficient option of the three since it minimises the number of task interactions), version `pi-v11.c` introduces the use of `taskgroup` and the `task_reduction` and `in_reduction` clauses. Take a look at the code and try to understand it: first, the `taskgroup` construct defines a task barrier at its end, so that the thread that is generating the tasks waits until they finish; second, the `taskgroup` construct also defines a reduction operation (`task_reduction(+: sum)` clause) that will occur at its end, gathering the contributions for variable `sum` from the 2 tasks generated; and third, each one of the tasks specifies that it will contribute to the reduction with its results for variable `sum` (with the `in_reduction(+: sum)` clause). Compile and run the `-debug` binary to check that the result is correct. Compile and run the `-omp` version to measure its scalability and check that it is close to the ideal case.
4. `pi-v12.c` provides an alternative version based on the use of `atomic` to protect the update of shared variable `sum`. In this case the code makes use of the `taskwait` construct to simply wait for the termination of the tasks: it forces the thread that entered into the `single` region and created the two tasks to wait for their termination before using the global value for `sum` that they have computed. Execute both the `-debug` and `-omp` binaries to validate the execution and notice again the large overhead introduced by the use of `atomic` (which would be much terrible if `critical` had been used).
5. `pi-v13.c` introduces the use of dependencies between tasks. The `depend` clause is used to specify variables that are `in`, `out` or `inout` to the task (in other words, read, written or both). With this specification the OpenMP runtime can establish the appropriate execution order for tasks so that any task is not executed until all its dependences are satisfied. This is the case for the third task defined in `pi-v13.c`, which waits for the termination of the other two. This is replacing the use of `taskgroup` and the reduction clauses in `pi-v11.c` or the use of `taskwait` and `atomic` (or `critical`) in `pi-v12.c`. Execute both the `-debug` and `-omp` binaries to validate the execution and appreciate the overheads of using task dependences.
6. Finally, `pi-v14.c` makes use of the `taskloop` construct to generate a task for a certain number of consecutive iterations, controlled either with the `num_tasks` or the `grainsize` clauses; the first one specifies the number of tasks to generate while the second one controls the number of consecutive iterations per task. Take a look at the code and try both options commenting one option or the other (and different values). For example, when `num_tasks(8)` is used, in the `-debug` version try to see which thread executes each task and how many tasks each thread executes; run several times if necessary. Compile and run the `-omp` version to measure its scalability and check that it is close to the ideal case.

2.1.2 Summary of code versions

The following table summarises all the codes used during the process. Changes accumulate from one version to the following.

Version	Description of changes	Correct?
v9	Use of <code>task</code> construct	no
v10	Use of <code>single</code> construct to have a single task generator	no
v11	Use of <code>taskgroup</code> and reductions (<code>task_reduction</code> and <code>in_reduction</code>)	yes
v12	Use of <code>taskwait</code> and <code>atomic</code>	yes
v13	Use of <code>task</code> with dependences (<code>depend</code> clause)	yes
v14	Use of <code>taskloop</code> to generate tasks from loop iterations	yes

2.2 Test your understanding

Next you will go through a set of simple examples that will be helpful to practice the components of the OpenMP programming model that have been used to parallelise the Pi computation. For each example answer the question(s) in the questionnaire that you will deliver as part of the deliverable for this second laboratory assignment. In order to follow them, you will need:

- The set of files inside the directory `H0_OpenMP/openmp/Part2`. Open each file in the directory (the examples are ordered), compile and execute; then answer the question(s) in the questionnaire associated to it. If asked in the program, do the required modifications.
- In case you want (or need) to know more details about the OpenMP constructs, we suggest you either use the complete specification document or the short reference guide that you will find either at the www.openmp.org website or at Atenea.

2.3 Observing overheads

To finish with this section we propose you to measure the overheads related with the creation of `parallel` regions, `task` creation and synchronisation. Codes are available inside directory `H0_OpenMP/overheads`. Take note of all the results that you obtain and reach your conclusions about the overheads associated with these OpenMP constructs. You will have to deliver them in the appropriate section of the deliverable for this second laboratory assignment.

2.3.1 Thread creation and termination

1. Open the `pi_omp_parallel.c` file and look at the changes done to the parallel version of `pi`. Function `difference` returns the difference in time between the sequential execution and the parallel execution when using a certain number of threads, using the OpenMP intrinsic function `omp_set_num_threads` to change the number of threads. The execution of both the sequential and parallel versions is done `NUMITERS` times in order to average the execution time of one iteration. In the `main` program, there is a loop to iterate over the number of threads (between 2 and `max_threads`, one of the input arguments of the execution) and prints the difference in time reported by function `difference`, which is the overhead introduced by the `parallel` construct (try to understand what is printed there).
2. Compile using the appropriate target in `Makefile` and submit the execution of the binary generated `pi_omp_parallel` with just one iteration and a maximum of 20 threads (i.e. `"sbatch ./submit-omp.sh pi_omp_parallel 1 20"`). Do not expect a correct result for the value of Pi!
3. How does the overhead of creating/terminating threads varies with the number of threads used? Which is the order of magnitude for the overhead of creating/terminating each individual thread in the parallel region?

2.3.2 Task creation and synchronisation

1. Open the `pi_omp_tasks.c` file. This version is creating tasks inside function `difference` by a single thread in the parallel region. The function measures the difference between the sequential execution and the version that creates the tasks; each version is repeated `NUMITERS` times in order to average the execution time of one iteration. In the `main` program, there is a loop to iterate over the number of tasks that we want to generate (between `MINTASKS` and `MAXTASKS`, and prints the difference in time reported by function `difference`, which is the overhead introduced by the `task` and `taskwait` constructs (try to understand what is printed there).
2. Compile using the appropriate target in `Makefile` and submit the execution of the binary generated `pi_omp_tasks` with just 10 iterations and one thread (i.e. `"sbatch ./submit-omp.sh pi_omp_tasks 10 1"`). Do not expect a correct result for the value of Pi!

3. How does the overhead of creating/synchronising tasks varies with the number of tasks created?
Which is the order of magnitude for the overhead of creating/synchronising each individual task?

Examples 3

A very practical introduction to OpenMP (Part III)

3.1 Work-sharing constructs

Work-sharing constructs divide the execution of code enclosed by the construct between threads in a team. For work-sharing to take place, the construct must be enclosed within the dynamic extent of a parallel region. A work-sharing construct distributes the execution of the corresponding region among the members of the team that encounters it. Threads execute portions of the region in the context of the implicit tasks that each one is executing. If the team consists of only one thread then the work-sharing region is not executed in parallel.

A work-sharing region has no barrier on entry; however, an *implicit* barrier exists at the end of the work-sharing region, unless a `nowait` clause is specified.

Directory `H0_OpenMP/openmp/worksharing` contain a set of codes related to *work sharing* constructs. Inspect the code and the behaviour of the parallel execution trying to learn from them.

3.2 Additional examples

In order to complete the set of examples we end up with a few additional examples which can be found in `H0_OpenMP/openmp/other`.